

StackOverflow

```
// Input file
z.put("stackoverflow", "/home/ubuntu/test_spark/stackoverflow/stackoverflow.csv")

%sh
head {stackoverflow}

1,27233496,,,0,C#
1,23698767,,,9,C#
1,5484340,,,0,C#
2,5494879,5484340,1,
1,9419744,,,2,Objective-C
1,26875732,,,1,C#
1,9002525,,,2,C++
2,9003401,9002525,4,
2,9003942,9002525,1,
2,9005311,9002525,0,

/** A raw stackoverflow posting, either a question or an answer */

type QID = Int
type HighScore = Int
type LangIndex = Int

case class Posting(postingType: Int, id: Int, acceptedAnswer: Option[Int], parentId: Option[QID], score: Int, tags: Option[String]) extends Serializable

type Question = Posting
type Answer = Posting

defined type alias QID
defined type alias HighScore
defined type alias LangIndex
defined class Posting
defined type alias Question
defined type alias Answer

// Start reading and transforming csv file

val lines = sc.textFile(z.get("stackoverflow").toString)

lines.take(5)

lines: org.apache.spark.rdd.RDD[String] = /home/ubuntu/test_spark/stackoverflow/stackoverflow.csv MapPartitionsRDD[697] at textFile at <console>:169
res159: Array[String] = Array(1,27233496,,,0,C#, 1,23698767,,,9,C#, 1,5484340,,,0,C#, 2,5494879,5484340,1,, 1,9419744,,,2,Objective-C)

/** Load postings from the given file */

/**
 * <postId>,<id>,<acceptedAnswer>,<parentId>,<score>,<tag>
 *
 * A short explanation of the comma-separated fields follows.
 * <postId>:      Type of the post. Type 1 = question, type 2 = answer.
 * <id>:          Unique id of the post (regardless of type).
 * <acceptedAnswer>: Id of the accepted answer post. This information is optional, so maybe be missingindicated by an empty string.
 * <parentId>:     For an answer: id of the corresponding question. For a question:missing, indicated by an empty string.
 * <score>:        The StackOverflow score (based on user votes).
 * <tag>:          The tag indicates the programming language that the post is about, in case it's a question, or missing in case it's an answer.
 */

import org.apache.spark.rdd.RDD

def rawPostings(lines: RDD[String]): RDD[Posting] = {
  lines.map(line => {
    val elems = line.split(",")
    Posting(postingType = elems(0).toInt,
            id = elems(1).toInt,
            acceptedAnswer = if (elems(2) == "") None else Some(elems(2).toInt),
            parentId = if (elems(3) == "") None else Some(elems(3).toInt),
            score = elems(4).toInt,
            tags = if (elems.length >= 6) Some(elems(5).intern()) else None)
  })
}

import org.apache.spark.rdd.RDD
rawPostings: (lines: org.apache.spark.rdd.RDD[String])org.apache.spark.rdd.RDD[Posting]

val raw = rawPostings(lines)
raw.take(5)

raw: org.apache.spark.rdd.RDD[Posting] = MapPartitionsRDD[698] at map at <console>:185
res160: Array[Posting] = Array(Posting(1,27233496,None,None,0,Some(C#)), Posting(1,23698767,None,None,9,Some(C#)), Posting(1,5484340,None,None,0,Some(C#)), Posting(2,5494879,None,Some(5484340),1,None), P

import org.apache.spark.HashPartitioner
val q = raw
  .filter(_.postingType == 1)
  .map(question => (question.id, question))
  .partitionBy(new HashPartitioner(100))
q.take(5)

import org.apache.spark.HashPartitioner
q: org.apache.spark.rdd.RDD[(Int, Posting)] = ShuffledRDD[701] at partitionBy at <console>:170
res161: Array[(Int, Posting)] = Array((3406000,Posting(1,3406000,None,None,0,Some(C#))), (17520100,Posting(1,17520100,None,None,1,Some(JavaScript))), (23620700,Posting(1,23620700,None,None,0,Some(PHP))),

q.partitions.length

res162: Int = 100

// ** repartition is not an action **

val q1 = q.repartition(10)

q1: org.apache.spark.rdd.RDD[(Int, Posting)] = MapPartitionsRDD[705] at repartition at <console>:165

val answers = raw
  .filter(_.postingType == 2)
  .map(answer => (answer.parentId.get, answer))
  .partitionBy(new HashPartitioner(100))
answers.take(5)

answers: org.apache.spark.rdd.RDD[(QID, Posting)] = ShuffledRDD[644] at partitionBy at <console>:154
res142: Array[(QID, Posting)] = Array((3406000,Posting(2,3406000,None,Some(3406000),0,None)), (3406000,Posting(2,3408739,None,Some(3406000),0,None)), (17933700,Posting(2,17935422,None,Some(17933700),5,No

answers.partitions.length

res84: Int = 100

val joined = q.join(answers)
joined.take(5)

joined: org.apache.spark.rdd.RDD[(Int, (Posting, Posting))] = MapPartitionsRDD[647] at join at <console>:153
res143: Array[(Int, (Posting, Posting))] = Array((7885200,(Posting(1,7885200,None,None,0,Some(JavaScript)),Posting(2,7885231,None,Some(7885200),2,None))), (637100,(Posting(1,637100,None,None,8,Some(Java)

val groupedPostings = joined.groupByKey()
groupedPostings.take(5)

groupedPostings: org.apache.spark.rdd.RDD[(Int, Iterable[(Posting, Posting)]] = MapPartitionsRDD[648] at groupByKey at <console>:151
res144: Array[(Int, Iterable[(Posting, Posting)]] = Array((7885200,CompactBuffer((Posting(1,7885200,None,None,0,Some(JavaScript)),Posting(2,7885231,None,Some(7885200),2,None)))), (20746400,CompactBuffer

/** Group the questions and answers together */

def groupedPostings(postings: RDD[Posting]): RDD[(QID, Iterable[(Question, Answer)])] = {
  // Create questions rdd (id, question)
  val questions = postings
    .filter(_.postingType == 1)
    .map(question => (question.id, question))
    .partitionBy(new HashPartitioner(100))
```

```

    // Create answers rdd (parentId, answer)
    val answers = postings
      .filter(_.postingType == 2)
      .map(answer => (answer.parentId.get, answer))
      .partitionBy(new HashPartitioner(100))

    // 1 question will have multiple answers
    // Join both the rdds (inner join)
    val joined = questions.join(answers)

    // Finally group the results
    joined.groupByKey()
  }

groupedPostings: (postings: org.apache.spark.rdd.RDD[Posting])org.apache.spark.rdd.RDD[(QID, Iterable[(Question, Answer)])]

val grouped = groupedPostings(raw)
grouped.take(5)

grouped: org.apache.spark.rdd.RDD[(QID, Iterable[(Question, Answer)])] = MapPartitionsRDD[658] at groupByKey at <console>:179
res145: Array[(QID, Iterable[(Question, Answer)])] = Array((7885200,CompactBuffer[(Posting(1,7885200,None,None,0,Some(JavaScript)),Posting(2,7885231,None,Some(7885200),2,None))), (20746400,CompactBuffer

// Compute score

// val scored: RDD[(Question, HighScore)] = ???

val score = grouped.map(case (qid, qu_ans) => {
  // val question = qu_ans.head match {
  //   case (q: Question, ans: Answer) => q
  // }
  // (question, 1)

  val scores = qu_ans.map(case (q, ans) => ans.score)

  (qu_ans.head._1, scores, scores.max)
})

score.take(5)

score: org.apache.spark.rdd.RDD[(Question, Iterable[Int], Int)] = MapPartitionsRDD[659] at map at <console>:155
res146: Array[(Question, Iterable[Int], Int)] = Array((Posting(1,7885200,None,None,0,Some(JavaScript)),List(2),2), (Posting(1,20746400,None,None,1,Some(C#)),List(2, 2, 1),2), (Posting(1,637100,None,None,

/** Compute the maximum score for each posting */
def scoredPostings(grouped: RDD[(QID, Iterable[(Question, Answer)])]): RDD[(Question, HighScore)] = {
  val score = grouped.map(case (qid, qu_ans) => {
    (qu_ans.head._1, qu_ans.map(case (q, ans) => ans.score).max)
  })
  score
}

scoredPostings: (grouped: org.apache.spark.rdd.RDD[(QID, Iterable[(Question, Answer)])])org.apache.spark.rdd.RDD[(Question, HighScore)]

val scored = scoredPostings(grouped)
scored.take(5)

scored: org.apache.spark.rdd.RDD[(Question, HighScore)] = MapPartitionsRDD[660] at map at <console>:160
res147: Array[(Question, HighScore)] = Array((Posting(1,7885200,None,None,0,Some(JavaScript)),2), (Posting(1,20746400,None,None,1,Some(C#)),2), (Posting(1,637100,None,None,8,Some(Java)),8), (Posting(1,15

/** Languages */
val langs = List(
  "JavaScript", "Java", "PHP", "Python", "C#", "C++", "Ruby", "CSS",
  "Objective-C", "Perl", "Scala", "Haskell", "MATLAB", "Clojure", "Groovy")

/** Compute the vectors for the kmeans */
def vectorPostings(scored: RDD[(Question, HighScore)]): RDD[(LangIndex, HighScore)] = {
  val vector = scored.map(case (q, highScore) => {
    val tag = q.tags.getOrElse("")
    val index = langs.indexOf(tag)
    (index * 50000, highScore)
  })
  vector
}

langs: List[String] = List(JavaScript, Java, PHP, Python, C#, C++, Ruby, CSS, Objective-C, Perl, Scala, Haskell, MATLAB, Clojure, Groovy)
vectorPostings: (scored: org.apache.spark.rdd.RDD[(Question, HighScore)])org.apache.spark.rdd.RDD[(LangIndex, HighScore)]

val vectors = vectorPostings(scored)
vectors.take(5)

vectors: org.apache.spark.rdd.RDD[(LangIndex, HighScore)] = MapPartitionsRDD[661] at map at <console>:164
res148: Array[(LangIndex, HighScore)] = Array((0,2), (200000,2), (50000,8), (50000,2), (50000,3))

assert(vectors.count() == 2121822, "Incorrect number of vectors: " + vectors.count())

vectors.count

res157: Long = 2121822

vectors.par

```